

## SparkSession and SparkContext

```
In [1]: from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("classification").getOrCreate()
sc = spark.sparkContext
```

## Schema

```
In [2]: from pyspark.sql.types import StructType, StructField, IntegerType, DoubleType
```

```
In [3]: colNames = ["Elevation", "Aspect", "Slope",
                    "Horizontal_Distance_To_Hydrology", "Vertical_Distance_To_Hydrology",
                    "Horizontal_Distance_To_Roadways",
                    "Hillshade_9am", "Hillshade_Noon", "Hillshade_3pm",
                    "Horizontal_Distance_To_Fire_Points"]

for i in range(4):
    colNames += ["Wilderness_Area_"+str(i),]

for i in range(40):
    colNames += ["Soil_Type_"+str(i),]

colNames += ["Cover_Type",]
```

```
In [4]: schema = StructType()
for name in colNames:
    if name == "Cover_Type":
        schema.add(StructField(name, DoubleType(), True))
    else:
        schema.add(StructField(name, IntegerType(), True))
```

## Read CSV with the prepared schema

```
In [5]: data = spark.read.csv('../Dropbox/pj_ss/covtype/covtype.data', header=False, schema=schema)
```

```
In [6]: data.printSchema()
```

```
root
|-- Elevation: integer (nullable = true)
|-- Aspect: integer (nullable = true)
|-- Slope: integer (nullable = true)
|-- Horizontal_Distance_To_Hydrology: integer (nullable = true)
|-- Vertical_Distance_To_Hydrology: integer (nullable = true)
|-- Horizontal_Distance_To_Roadways: integer (nullable = true)
|-- Hillshade_9am: integer (nullable = true)
|-- Hillshade_Noon: integer (nullable = true)
|-- Hillshade_3pm: integer (nullable = true)
|-- Horizontal_Distance_To_Fire_Points: integer (nullable = true)
|-- Wilderness_Area_0: integer (nullable = true)
|-- Wilderness_Area_1: integer (nullable = true)
|-- Wilderness_Area_2: integer (nullable = true)
|-- Wilderness_Area_3: integer (nullable = true)
|-- Soil_Type_0: integer (nullable = true)
|-- Soil_Type_1: integer (nullable = true)
|-- Soil_Type_2: integer (nullable = true)
|-- Soil_Type_3: integer (nullable = true)
|-- Soil_Type_4: integer (nullable = true)
|-- Soil_Type_5: integer (nullable = true)
|-- Soil_Type_6: integer (nullable = true)
|-- Soil_Type_7: integer (nullable = true)
|-- Soil_Type_8: integer (nullable = true)
|-- Soil_Type_9: integer (nullable = true)
|-- Soil_Type_10: integer (nullable = true)
|-- Soil_Type_11: integer (nullable = true)
|-- Soil_Type_12: integer (nullable = true)
|-- Soil_Type_13: integer (nullable = true)
|-- Soil_Type_14: integer (nullable = true)
|-- Soil_Type_15: integer (nullable = true)
|-- Soil_Type_16: integer (nullable = true)
|-- Soil_Type_17: integer (nullable = true)
|-- Soil_Type_18: integer (nullable = true)
|-- Soil_Type_19: integer (nullable = true)
|-- Soil_Type_20: integer (nullable = true)
|-- Soil_Type_21: integer (nullable = true)
|-- Soil_Type_22: integer (nullable = true)
|-- Soil_Type_23: integer (nullable = true)
|-- Soil_Type_24: integer (nullable = true)
|-- Soil_Type_25: integer (nullable = true)
|-- Soil_Type_26: integer (nullable = true)
|-- Soil_Type_27: integer (nullable = true)
|-- Soil_Type_28: integer (nullable = true)
|-- Soil_Type_29: integer (nullable = true)
|-- Soil_Type_30: integer (nullable = true)
|-- Soil_Type_31: integer (nullable = true)
|-- Soil_Type_32: integer (nullable = true)
|-- Soil_Type_33: integer (nullable = true)
|-- Soil_Type_34: integer (nullable = true)
|-- Soil_Type_35: integer (nullable = true)
|-- Soil_Type_36: integer (nullable = true)
|-- Soil_Type_37: integer (nullable = true)
|-- Soil_Type_38: integer (nullable = true)
|-- Soil_Type_39: integer (nullable = true)
|-- Cover_Type: double (nullable = true)
```

```
In [7]: data.take(1)
```

```
Out[7]: [Row(Elevation=2596, Aspect=51, Slope=3, Horizontal_Distance_To_Hydrology=258, Vertical_Distance_To_Hydrology=0, Horizontal_Distance_To_Roadways=510, Hillshade_9am=221, Hillshade_Noon=232, Hillshade_3pm=148, Horizontal_Distance_To_Fire_Points=6279, Wilderness_Area_0=1, Wilderness_Area_1=0, Wilderness_Area_2=0, Wilderness_Area_3=0, Soil_Type_0=0, Soil_Type_1=0, Soil_Type_2=0, Soil_Type_3=0, Soil_Type_4=0, Soil_Type_5=0, Soil_Type_6=0, Soil_Type_7=0, Soil_Type_8=0, Soil_Type_9=0, Soil_Type_10=0, Soil_Type_11=0, Soil_Type_12=0, Soil_Type_13=0, Soil_Type_14=0, Soil_Type_15=0, Soil_Type_16=0, Soil_Type_17=0, Soil_Type_18=0, Soil_Type_19=0, Soil_Type_20=0, Soil_Type_21=0, Soil_Type_22=0, Soil_Type_23=0, Soil_Type_24=0, Soil_Type_25=0, Soil_Type_26=0, Soil_Type_27=0, Soil_Type_28=1, Soil_Type_29=0, Soil_Type_30=0, Soil_Type_31=0, Soil_Type_32=0, Soil_Type_33=0, Soil_Type_34=0, Soil_Type_35=0, Soil_Type_36=0, Soil_Type_37=0, Soil_Type_38=0, Soil_Type_39=0, Cover_Type=5.0)]
```

## Assembler

```
In [8]: from pyspark.ml.linalg import Vectors
        from pyspark.ml.feature import VectorAssembler
```

```
In [9]: (trainData, testData) = data.randomSplit([0.9, 0.1])
```

```
In [10]: trainData
```

```
Out[10]: DataFrame[Elevation: int, Aspect: int, Slope: int, Horizontal_Distance_To_Hydrology: int, Vertical_Distance_To_Hydrology: int, Horizontal_Distance_To_Roadways: int, Hillshade_9am: int, Hillshade_Noon: int, Hillshade_3pm: int, Horizontal_Distance_To_Fire_Points: int, Wilderness_Area_0: int, Wilderness_Area_1: int, Wilderness_Area_2: int, Wilderness_Area_3: int, Soil_Type_0: int, Soil_Type_1: int, Soil_Type_2: int, Soil_Type_3: int, Soil_Type_4: int, Soil_Type_5: int, Soil_Type_6: int, Soil_Type_7: int, Soil_Type_8: int, Soil_Type_9: int, Soil_Type_10: int, Soil_Type_11: int, Soil_Type_12: int, Soil_Type_13: int, Soil_Type_14: int, Soil_Type_15: int, Soil_Type_16: int, Soil_Type_17: int, Soil_Type_18: int, Soil_Type_19: int, Soil_Type_20: int, Soil_Type_21: int, Soil_Type_22: int, Soil_Type_23: int, Soil_Type_24: int, Soil_Type_25: int, Soil_Type_26: int, Soil_Type_27: int, Soil_Type_28: int, Soil_Type_29: int, Soil_Type_30: int, Soil_Type_31: int, Soil_Type_32: int, Soil_Type_33: int, Soil_Type_34: int, Soil_Type_35: int, Soil_Type_36: int, Soil_Type_37: int, Soil_Type_38: int, Soil_Type_39: int, Cover_Type: double]
```

```
In [11]: inputCols = trainData.drop('Cover_Type').columns
```

```
In [12]: inputCols
```

```
Out[12]: ['Elevation',
          'Aspect',
          'Slope',
          'Horizontal_Distance_To_Hydrology',
          'Vertical_Distance_To_Hydrology',
          'Horizontal_Distance_To_Roadways',
          'Hillshade_9am',
          'Hillshade_Noon',
          'Hillshade_3pm',
          'Horizontal_Distance_To_Fire_Points',
          'Wilderness_Area_0',
          'Wilderness_Area_1',
          'Wilderness_Area_2',
          'Wilderness_Area_3',
          'Soil_Type_0',
          'Soil_Type_1',
          'Soil_Type_2',
          'Soil_Type_3',
          'Soil_Type_4',
          'Soil_Type_5',
          'Soil_Type_6',
          'Soil_Type_7',
          'Soil_Type_8',
          'Soil_Type_9',
          'Soil_Type_10',
          'Soil_Type_11',
          'Soil_Type_12',
          'Soil_Type_13',
          'Soil_Type_14',
          'Soil_Type_15',
          'Soil_Type_16',
          'Soil_Type_17',
          'Soil_Type_18',
          'Soil_Type_19',
          'Soil_Type_20',
          'Soil_Type_21',
          'Soil_Type_22',
          'Soil_Type_23',
          'Soil_Type_24',
          'Soil_Type_25',
          'Soil_Type_26',
          'Soil_Type_27',
          'Soil_Type_28',
          'Soil_Type_29',
          'Soil_Type_30',
          'Soil_Type_31',
          'Soil_Type_32',
          'Soil_Type_33',
          'Soil_Type_34',
          'Soil_Type_35',
          'Soil_Type_36',
          'Soil_Type_37',
          'Soil_Type_38',
          'Soil_Type_39']
```

```
In [16]: assembler = VectorAssembler(
          inputCols=inputCols,
          outputCol="featureVector")
```

```
In [17]: assembler
```

```
Out[17]: VectorAssembler_4bb3bb34a9cf0b6fd419
```

```
In [18]: assembledTrainData = assembler.transform(trainData)
```

```
In [19]: assembledTrainData
```

```
Out[19]: DataFrame[Elevation: int, Aspect: int, Slope: int, Horizontal_Distance_
To_Hydrology: int, Vertical_Distance_To_Hydrology: int, Horizontal_Distance_To_Roadways: int, Hillshade_9am: int, Hillshade_Noon: int, Hillshade_3pm: int, Horizontal_Distance_To_Fire_Points: int, Wilderness_Area_0: int, Wilderness_Area_1: int, Wilderness_Area_2: int, Wilderness_Area_3: int, Soil_Type_0: int, Soil_Type_1: int, Soil_Type_2: int, Soil_Type_3: int, Soil_Type_4: int, Soil_Type_5: int, Soil_Type_6: int, Soil_Type_7: int, Soil_Type_8: int, Soil_Type_9: int, Soil_Type_10: int, Soil_Type_11: int, Soil_Type_12: int, Soil_Type_13: int, Soil_Type_14: int, Soil_Type_15: int, Soil_Type_16: int, Soil_Type_17: int, Soil_Type_18: int, Soil_Type_19: int, Soil_Type_20: int, Soil_Type_21: int, Soil_Type_22: int, Soil_Type_23: int, Soil_Type_24: int, Soil_Type_25: int, Soil_Type_26: int, Soil_Type_27: int, Soil_Type_28: int, Soil_Type_29: int, Soil_Type_30: int, Soil_Type_31: int, Soil_Type_32: int, Soil_Type_33: int, Soil_Type_34: int, Soil_Type_35: int, Soil_Type_36: int, Soil_Type_37: int, Soil_Type_38: int, Soil_Type_39: int, Cover_Type: double, featureVector: vector]
```

```
In [20]: assembledTrainData.select('featureVector').show(3, truncate=False)
```

```
+-----+
+-----+
|featureVector|
|             |
+-----+
+-----+
|(54,[0,1,2,3,4,5,6,7,8,9,13,15],[1863.0,37.0,17.0,120.0,18.0,90.0,217.0,202.0,115.0,769.0,1.0,1.0])|
|(54,[0,1,2,3,4,5,6,7,8,9,13,18],[1879.0,28.0,19.0,30.0,12.0,95.0,209.0,196.0,117.0,778.0,1.0,1.0])|
|(54,[0,1,2,3,4,5,6,7,8,9,13,15],[1888.0,33.0,22.0,150.0,46.0,108.0,209.0,185.0,103.0,735.0,1.0,1.0])|
+-----+
+-----+
only showing top 3 rows
```

## Decision Tree

```
In [21]: from pyspark.ml.classification import DecisionTreeClassifier
```

```
In [22]: classifier = DecisionTreeClassifier(labelCol="Cover_Type",  
                                           featuresCol="featureVector",  
                                           predictionCol="prediction")
```

```
In [23]: model = classifier.fit(assembledTrainData)
```

```
In [24]: model
```

```
Out[24]: DecisionTreeClassificationModel (uid=DecisionTreeClassifier_44ac953bdb2  
a9ad6e065) of depth 5 with 63 nodes
```

```
In [25]: print(model.toDebugString)
```



DecisionTreeClassificationModel (uid=DecisionTreeClassifier\_44ac953bdb2a9ad6e065) of depth 5 with 63 nodes

```
If (feature 0 <= 3034.5)
  If (feature 0 <= 2479.5)
    If (feature 3 <= 15.0)
      If (feature 13 <= 0.5)
        If (feature 17 <= 0.5)
          Predict: 6.0
        Else (feature 17 > 0.5)
          Predict: 6.0
      Else (feature 13 > 0.5)
        If (feature 23 <= 0.5)
          Predict: 4.0
        Else (feature 23 > 0.5)
          Predict: 3.0
    Else (feature 3 > 15.0)
      If (feature 16 <= 0.5)
        If (feature 9 <= 576.5)
          Predict: 3.0
        Else (feature 9 > 576.5)
          Predict: 3.0
      Else (feature 16 > 0.5)
        If (feature 9 <= 1295.0)
          Predict: 3.0
        Else (feature 9 > 1295.0)
          Predict: 4.0
  Else (feature 0 > 2479.5)
    If (feature 17 <= 0.5)
      If (feature 15 <= 0.5)
        If (feature 0 <= 2942.5)
          Predict: 2.0
        Else (feature 0 > 2942.5)
          Predict: 2.0
      Else (feature 15 > 0.5)
        If (feature 9 <= 1380.5)
          Predict: 3.0
        Else (feature 9 > 1380.5)
          Predict: 3.0
    Else (feature 17 > 0.5)
      If (feature 0 <= 2692.5)
        If (feature 0 <= 2635.5)
          Predict: 3.0
        Else (feature 0 > 2635.5)
          Predict: 3.0
      Else (feature 0 > 2692.5)
        If (feature 5 <= 1200.5)
          Predict: 5.0
        Else (feature 5 > 1200.5)
          Predict: 2.0
  Else (feature 0 > 3034.5)
    If (feature 0 <= 3309.5)
      If (feature 7 <= 238.5)
        If (feature 0 <= 3101.5)
          If (feature 3 <= 191.0)
            Predict: 1.0
          Else (feature 3 > 191.0)
            Predict: 2.0
```

```

Else (feature 0 > 3101.5)
  If (feature 5 <= 998.0)
    Predict: 1.0
  Else (feature 5 > 998.0)
    Predict: 1.0
Else (feature 7 > 238.5)
  If (feature 3 <= 333.0)
    If (feature 0 <= 3186.5)
      Predict: 1.0
    Else (feature 0 > 3186.5)
      Predict: 1.0
  Else (feature 3 > 333.0)
    If (feature 0 <= 3206.5)
      Predict: 2.0
    Else (feature 0 > 3206.5)
      Predict: 1.0
Else (feature 0 > 3309.5)
  If (feature 12 <= 0.5)
    If (feature 3 <= 290.0)
      If (feature 6 <= 207.5)
        Predict: 1.0
      Else (feature 6 > 207.5)
        Predict: 7.0
    Else (feature 3 > 290.0)
      If (feature 10 <= 0.5)
        Predict: 1.0
      Else (feature 10 > 0.5)
        Predict: 1.0
  Else (feature 12 > 0.5)
    If (feature 45 <= 0.5)
      If (feature 0 <= 3369.5)
        Predict: 7.0
      Else (feature 0 > 3369.5)
        Predict: 7.0
    Else (feature 45 > 0.5)
      If (feature 5 <= 998.0)
        Predict: 7.0
      Else (feature 5 > 998.0)
        Predict: 1.0

```

```
In [26]: print(model.featureImportances)
```

```

(54,[0,3,5,6,7,9,10,12,13,15,16,17,23,45],[0.8133742673793523,0.0309786
0477331104,0.014747043247755682,0.0023999635668687604,0.025523595174575
104,0.00646850931723158,0.0031668991984998784,0.011192823508553537,0.00
27142714212239146,0.03035888049891708,0.0033263230701791355,0.038074558
83308248,0.0009625115515200479,0.01671174845892951])

```

```
In [27]: predictions = model.transform(assembledTrainData)
```

```
In [28]: predictions.select(["Cover_Type", "prediction", "probability"]).show(3,
truncate=False)
```

```
+-----+-----+-----+
|Cover_Type|prediction|probability|
+-----+-----+-----+
|6.0       |3.0       |[0.0,5.213492518638236E-5,0.0570356081539023,0.5660288827485532,0.025024764089463532,0.0,0.35185861008289454,0.0]|
|6.0       |3.0       |[0.0,5.213492518638236E-5,0.0570356081539023,0.5660288827485532,0.025024764089463532,0.0,0.35185861008289454,0.0]|
|6.0       |3.0       |[0.0,5.213492518638236E-5,0.0570356081539023,0.5660288827485532,0.025024764089463532,0.0,0.35185861008289454,0.0]|
+-----+-----+-----+
only showing top 3 rows
```

## Evaluation

```
In [29]: from pyspark.ml.evaluation import MulticlassClassificationEvaluator
```

```
In [30]: evaluator = MulticlassClassificationEvaluator(labelCol="Cover_Type",
predictionCol="prediction"
)
```

```
In [31]: evaluator
```

```
Out[31]: MulticlassClassificationEvaluator_4a96ab9498d109c86b86
```

```
In [32]: evaluator.setMetricName("accuracy").evaluate(predictions)
```

```
Out[32]: 0.7022924849787542
```

```
In [33]: evaluator.setMetricName("f1").evaluate(predictions)
```

```
Out[33]: 0.6857631073030778
```

## Pipeline

```
In [34]: from pyspark.ml import Pipeline
```

```
In [35]: inputCols = trainData.columns[:-1]
```

```
In [36]: inputCols
```

```
Out[36]: ['Elevation',
          'Aspect',
          'Slope',
          'Horizontal_Distance_To_Hydrology',
          'Vertical_Distance_To_Hydrology',
          'Horizontal_Distance_To_Roadways',
          'Hillshade_9am',
          'Hillshade_Noon',
          'Hillshade_3pm',
          'Horizontal_Distance_To_Fire_Points',
          'Wilderness_Area_0',
          'Wilderness_Area_1',
          'Wilderness_Area_2',
          'Wilderness_Area_3',
          'Soil_Type_0',
          'Soil_Type_1',
          'Soil_Type_2',
          'Soil_Type_3',
          'Soil_Type_4',
          'Soil_Type_5',
          'Soil_Type_6',
          'Soil_Type_7',
          'Soil_Type_8',
          'Soil_Type_9',
          'Soil_Type_10',
          'Soil_Type_11',
          'Soil_Type_12',
          'Soil_Type_13',
          'Soil_Type_14',
          'Soil_Type_15',
          'Soil_Type_16',
          'Soil_Type_17',
          'Soil_Type_18',
          'Soil_Type_19',
          'Soil_Type_20',
          'Soil_Type_21',
          'Soil_Type_22',
          'Soil_Type_23',
          'Soil_Type_24',
          'Soil_Type_25',
          'Soil_Type_26',
          'Soil_Type_27',
          'Soil_Type_28',
          'Soil_Type_29',
          'Soil_Type_30',
          'Soil_Type_31',
          'Soil_Type_32',
          'Soil_Type_33',
          'Soil_Type_34',
          'Soil_Type_35',
          'Soil_Type_36',
          'Soil_Type_37',
          'Soil_Type_38',
          'Soil_Type_39']
```

```
In [37]: assembler = VectorAssembler(inputCols=inputCols, outputCol="featureVector")
```

```
In [38]: classifier = DecisionTreeClassifier(labelCol="Cover_Type",  
                                             featuresCol="featureVector",  
                                             predictionCol="prediction")
```

```
In [39]: classifier
```

```
Out[39]: DecisionTreeClassifier_41c58c49ca4ad38a5bc6
```

```
In [40]: pipeline = Pipeline(stages=[assembler, classifier])
```

```
In [41]: pipeline
```

```
Out[41]: Pipeline_4b6aad06f5c5ac75d01b
```

```
In [42]: from pyspark.ml.tuning import ParamGridBuilder
```

```
In [46]: paramGrid = ParamGridBuilder()\  
        .addGrid(classifier, ["gini", "entropy"])\  
        .addGrid(classifier, [1, 20])\  
        .addGrid(classifier, [40, 300])\  
        .addGrid(classifier, [0.0, 0.05])\  
        .build()
```

```
In [47]: multiclassEval = MulticlassClassificationEvaluator(  
        labelCol="Cover_Type",  
        predictionCol="prediction",  
        metricName="accuracy")
```

```
In [48]: multiclassEval.evaluate(predictions)
```

```
Out[48]: 0.7022924849787542
```

```
In [49]: from pyspark.ml.tuning import TrainValidationSplit
```

```
In [50]: validator = TrainValidationSplit(  
        estimator=pipeline,  
        estimatorParamMaps=paramGrid,  
        evaluator=multiclassEval,  
        trainRatio=0.9)
```

```
In [51]: validatorModel = validator.fit(trainData)
```

```
In [52]: bestModel = validatorModel.bestModel
```

```
In [53]: bestModel
```

```
Out[53]: PipelineModel_44b3ad4e4f4671d6a89b
```

```
In [54]: bestModel.stages[-1].extractParamMap()
```

```
Out[54]: {Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='cacheNodeIds', doc='If false, the algorithm will pass trees to executors to match instances with nodes. If true, the algorithm will cache node IDs for each instance. Caching can speed up training of deeper trees.'): False,
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='checkpointInterval', doc='set checkpoint interval (>= 1) or disable checkpoint (-1). E.g. 10 means that the cache will get checkpointed every 10 iterations. Note: this setting will be ignored if the checkpoint directory is not set in the SparkContext'): 10,
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='featuresCol', doc='features column name'): 'featureVector',
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='impurity', doc='Criterion used for information gain calculation (case-insensitive). Supported options: entropy, gini'): 'gini',
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='labelCol', doc='label column name'): 'Cover_Type',
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='maxBins', doc='Max number of bins for discretizing continuous features. Must be >=2 and >= number of categories for any categorical feature.'): 32,
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='maxDepth', doc='Maximum depth of the tree. (>= 0) E.g., depth 0 means 1 leaf node; depth 1 means 1 internal node + 2 leaf nodes.'): 5,
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='memoryInMB', doc='Maximum memory in MB allocated to histogram aggregation.'): 256,
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='minInfoGain', doc='Minimum information gain for a split to be considered at a tree node.'): 0.0,
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='minInstancesPerNode', doc='Minimum number of instances each child must have after split. If a split causes the left or right child to have fewer than minInstancesPerNode, the split will be discarded as invalid. Should be >= 1.'): 1,
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='predictionCol', doc='prediction column name'): 'prediction',
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='probabilityCol', doc='Column name for predicted class conditional probabilities. Note: Not all models output well-calibrated probability estimates! These probabilities should be treated as confidences, not precise probabilities.'): 'probability',
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='rawPredictionCol', doc='raw prediction (a.k.a. confidence) column name'): 'rawPrediction',
  Param(parent='DecisionTreeClassifier_41c58c49ca4ad38a5bc6', name='seed', doc='random seed'): 7750783964360596788}
```

```
In [55]: paramsAndMetrics = validatorModel.validationMetrics
```

```
In [56]: paramsAndMetrics
```

```
Out[56]: [0.7012882447665056, 0.7012882447665056]
```

```
In [57]: multiclassEval.evaluate(bestModel.transform(testData))
```

```
Out[57]: 0.7040250662810316
```

## Undoing Onehot Coding

```
In [58]: wildernessCols = []  
        for i in range(4):  
            wildernessCols += ["Wilderness_Area_"+str(i),]
```

```
In [63]: wildernessAssembler = VectorAssembler(  
        inputCols=wildernessCols,  
        outputCol="wilderness")
```

```
In [64]: from pyspark.sql.functions import udf  
        from pyspark.sql.types import ArrayType, DoubleType, StructType
```

```
In [65]: unhotudf = udf(lambda x: float(x.toArray().nonzero()[0]), DoubleType())
```

```
In [66]: withWilderness = wildernessAssembler.transform(data)
```

```
In [67]: withWilderness.take(1)
```

```
Out[67]: [Row(Elevation=2596, Aspect=51, Slope=3, Horizontal_Distance_To_Hydrology=258, Vertical_Distance_To_Hydrology=0, Horizontal_Distance_To_Roadways=510, Hillshade_9am=221, Hillshade_Noon=232, Hillshade_3pm=148, Horizontal_Distance_To_Fire_Points=6279, Wilderness_Area_0=1, Wilderness_Area_1=0, Wilderness_Area_2=0, Wilderness_Area_3=0, Soil_Type_0=0, Soil_Type_1=0, Soil_Type_2=0, Soil_Type_3=0, Soil_Type_4=0, Soil_Type_5=0, Soil_Type_6=0, Soil_Type_7=0, Soil_Type_8=0, Soil_Type_9=0, Soil_Type_10=0, Soil_Type_11=0, Soil_Type_12=0, Soil_Type_13=0, Soil_Type_14=0, Soil_Type_15=0, Soil_Type_16=0, Soil_Type_17=0, Soil_Type_18=0, Soil_Type_19=0, Soil_Type_20=0, Soil_Type_21=0, Soil_Type_22=0, Soil_Type_23=0, Soil_Type_24=0, Soil_Type_25=0, Soil_Type_26=0, Soil_Type_27=0, Soil_Type_28=1, Soil_Type_29=0, Soil_Type_30=0, Soil_Type_31=0, Soil_Type_32=0, Soil_Type_33=0, Soil_Type_34=0, Soil_Type_35=0, Soil_Type_36=0, Soil_Type_37=0, Soil_Type_38=0, Soil_Type_39=0, Cover_Type=5.0, wilderness=SparseVector(4, {0: 1.0}))]
```

```
In [68]: withWilderness = withWilderness\  
        .drop(*wildernessCols)\  
        .withColumn("wilderness", unhotudf(withWilderness['wilderness']))
```

```
In [69]: withWilderness.take(1)
```

```
Out[69]: [Row(Elevation=2596, Aspect=51, Slope=3, Horizontal_Distance_To_Hydrology=258, Vertical_Distance_To_Hydrology=0, Horizontal_Distance_To_Roadways=510, Hillshade_9am=221, Hillshade_Noon=232, Hillshade_3pm=148, Horizontal_Distance_To_Fire_Points=6279, Soil_Type_0=0, Soil_Type_1=0, Soil_Type_2=0, Soil_Type_3=0, Soil_Type_4=0, Soil_Type_5=0, Soil_Type_6=0, Soil_Type_7=0, Soil_Type_8=0, Soil_Type_9=0, Soil_Type_10=0, Soil_Type_11=0, Soil_Type_12=0, Soil_Type_13=0, Soil_Type_14=0, Soil_Type_15=0, Soil_Type_16=0, Soil_Type_17=0, Soil_Type_18=0, Soil_Type_19=0, Soil_Type_20=0, Soil_Type_21=0, Soil_Type_22=0, Soil_Type_23=0, Soil_Type_24=0, Soil_Type_25=0, Soil_Type_26=0, Soil_Type_27=0, Soil_Type_28=1, Soil_Type_29=0, Soil_Type_30=0, Soil_Type_31=0, Soil_Type_32=0, Soil_Type_33=0, Soil_Type_34=0, Soil_Type_35=0, Soil_Type_36=0, Soil_Type_37=0, Soil_Type_38=0, Soil_Type_39=0, Cover_Type=5.0, wilderness=0.0)]
```

```
In [70]: soilCols = []  
        for i in range(40):  
            soilCols += ["Soil_Type_"+str(i),]
```

```
In [71]: soilAssembler = VectorAssembler(  
        inputCols=soilCols,  
        outputCol="soil")
```

```
In [72]: withWilderness = soilAssembler.transform(withWilderness)
```

```
In [73]: unencodedData = withWilderness\  
        .drop(*soilCols)\  
        .withColumn("soil", unhotudf(withWilderness['soil']))
```

```
In [74]: unencodedData.take(2)
```

```
Out[74]: [Row(Elevation=2596, Aspect=51, Slope=3, Horizontal_Distance_To_Hydrology=258, Vertical_Distance_To_Hydrology=0, Horizontal_Distance_To_Roadways=510, Hillshade_9am=221, Hillshade_Noon=232, Hillshade_3pm=148, Horizontal_Distance_To_Fire_Points=6279, Cover_Type=5.0, wilderness=0.0, soil=28.0),  
        Row(Elevation=2590, Aspect=56, Slope=2, Horizontal_Distance_To_Hydrology=212, Vertical_Distance_To_Hydrology=-6, Horizontal_Distance_To_Roadways=390, Hillshade_9am=220, Hillshade_Noon=235, Hillshade_3pm=151, Horizontal_Distance_To_Fire_Points=6225, Cover_Type=5.0, wilderness=0.0, soil=28.0)]
```

## Decision Tree with Unencoded Data

```
In [75]: (unencTrainData, unencTestData) = unencodedData.randomSplit([0.9, 0.1])
```

```
In [76]: from pyspark.ml.feature import VectorIndexer
```

```
In [77]: inputCols = unencTrainData.drop('Cover_Type').columns
```



```
In [78]: inputCols
```

```
Out[78]: ['Elevation',  
         'Aspect',  
         'Slope',  
         'Horizontal_Distance_To_Hydrology',  
         'Vertical_Distance_To_Hydrology',  
         'Horizontal_Distance_To_Roadways',  
         'Hillshade_9am',  
         'Hillshade_Noon',  
         'Hillshade_3pm',  
         'Horizontal_Distance_To_Fire_Points',  
         'wilderness',  
         'soil']
```

```
In [79]: assembler = VectorAssembler(  
         inputCols=inputCols,  
         outputCol="featureVector")
```

```
In [80]: indexer = VectorIndexer(  
         maxCategories=40,  
         inputCol="featureVector",  
         outputCol="indexedVector")
```

```
In [81]: classifier = DecisionTreeClassifier(  
         seed=42,  
         labelCol="Cover_Type",  
         featuresCol="indexedVector",  
         predictionCol="prediction")
```

```
In [82]: pipeline = Pipeline(stages=[assembler, indexer, classifier])
```

## Random Forrest Classifier

```
In [83]: from pyspark.ml.classification import RandomForestClassifier
```

```
In [84]: classifier = RandomForestClassifier(  
         seed=42,  
         maxBins=40,  
         labelCol="Cover_Type",  
         featuresCol="indexedVector",  
         predictionCol="prediction")
```

```
In [85]: pipeline = Pipeline(stages=[assembler, indexer, classifier])
```

```
In [86]: paramGrid = ParamGridBuilder()\  
         .addGrid(classifier.minInfoGain, [0.0, 0.05])\  
         .addGrid(classifier.numTrees, [1, 10])\  
         .build()
```

```
In [87]: multiclassEval = MulticlassClassificationEvaluator(  
        labelCol="Cover_Type",  
        predictionCol="prediction",  
        metricName="accuracy")
```

```
In [88]: validator = TrainValidationSplit(  
        seed=42,  
        estimator=pipeline,  
        evaluator=multiclassEval,  
        estimatorParamMaps=paramGrid,  
        trainRatio=0.9)
```

```
In [89]: validatorModel = validator.fit(unencTrainData)
```

```
In [90]: bestModel = validatorModel.bestModel
```

```
In [91]: bestModel
```

```
Out[91]: PipelineModel_4cdf8601e7c081412e4a
```

```
In [92]: forestModel = bestModel.stages[-1]
print(forestModel.extractParamMap())
```

```
{Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='cacheNodeIds', doc='If false, the algorithm will pass trees to executors to match instances with nodes. If true, the algorithm will cache node IDs for each instance. Caching can speed up training of deeper trees.'): False, Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='checkpointInterval', doc='set checkpoint interval (>= 1) or disable checkpoint (-1). E.g. 10 means that the cache will get checkpointed every 10 iterations. Note: this setting will be ignored if the checkpoint directory is not set in the SparkContext'): 10, Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='featureSubsetStrategy', doc='The number of features to consider for splits at each tree node. Supported options: auto, all, onethird, sqrt, log2, (0.0-1.0], [1-n].'): 'auto', Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='featuresCol', doc='features column name'): 'indexedVector', Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='impurity', doc='Criterion used for information gain calculation (case-insensitive). Supported options: entropy, gini'): 'gini', Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='labelCol', doc='label column name'): 'Cover_Type', Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='maxBins', doc='Max number of bins for discretizing continuous features. Must be >=2 and >= number of categories for any categorical feature.'): 40, Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='maxDepth', doc='Maximum depth of the tree. (>= 0) E.g., depth 0 means 1 leaf node; depth 1 means 1 internal node + 2 leaf nodes.'): 5, Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='maxMemoryInMB', doc='Maximum memory in MB allocated to histogram aggregation.'): 256, Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='minInfoGain', doc='Minimum information gain for a split to be considered at a tree node.'): 0.0, Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='minInstancesPerNode', doc='Minimum number of instances each child must have after split. If a split causes the left or right child to have fewer than minInstancesPerNode, the split will be discarded as invalid. Should be >= 1.'): 1, Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='numTrees', doc='Number of trees to train (>= 1)'): 1, Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='predictionCol', doc='prediction column name'): 'prediction', Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='probabilityCol', doc='Column name for predicted class conditional probabilities. Note: Not all models output well-calibrated probability estimates! These probabilities should be treated as confidence scores, not precise probabilities'): 'probability', Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='rawPredictionCol', doc='raw prediction (a.k.a. confidence) column name'): 'rawPrediction', Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='seed', doc='random seed'): 42, Param(parent='RandomForestClassifier_4fa2ae5257b23789060c', name='subsamplingRate', doc='Fraction of the training data used for learning each decision tree, in range (0, 1].'): 1.0}
```

```
In [93]: forestModel.getNumTrees
```

```
Out[93]: 1
```

```
In [94]: sorted(list(zip(inputCols, forestModel.featureImportances)), key=lambda  
x: x[1], reverse=True)
```

```
Out[94]: [('Elevation', 0.7822498415213323),  
( 'soil', 0.11297264686002277),  
( 'wilderness', 0.037105461899423466),  
( 'Horizontal_Distance_To_Roadways', 0.03535837398448177),  
( 'Horizontal_Distance_To_Fire_Points', 0.012732741797356662),  
( 'Hillshade_Noon', 0.011420066167183551),  
( 'Horizontal_Distance_To_Hydrology', 0.0050684199555167024),  
( 'Vertical_Distance_To_Hydrology', 0.003092447814683044),  
( 'Aspect', 0.0),  
( 'Slope', 0.0),  
( 'Hillshade_9am', 0.0),  
( 'Hillshade_3pm', 0.0)]
```

```
In [95]: testAccuracy = multiclassEval.evaluate(bestModel.transform(unencTestData  
) )
```

```
In [96]: testAccuracy
```

```
Out[96]: 0.6989563173901023
```